

PHASE 4 ASSIGNMENT

PROJECT TITLE: Feature selection, Model training, Evaluation of an dataset.

PROBLEM DEFINITION: The problem is to predict house prices using machine learning techniques. The objective is to develop a model that accurately predicts the prices of houses based on a set of features such as location, square footage, number of bedrooms and bathrooms, and other relevant factors. This project involves data preprocessing, feature engineering, model selection, training, and evaluation.

GITHUB LINK:

<https://github.com/JaisriS/predicting-house-prices-using-machine-learning.git>

<https://github.com/JaisriS/innovation.git>

DOCUMENT:

Building the project by Feature selection, Model training, Evaluation of an dataset.

DATASET LINK ON: Predicting House Prices

<https://www.kaggle.com/datasets/vedavyasv/usa-housing>

Creating a house price prediction model involves several key steps, including feature selection, model training, and evaluation. Here's a step-by-step guide to help you build such a model:

1. Data Collection and Preparation:

- Gather a dataset that includes information about houses and their sale prices. Common features might include square footage, number of bedrooms and bathrooms, location, etc.

- Preprocess the data by handling missing values, encoding categorical variables (e.g., one-hot encoding), and scaling numerical features if necessary.

2. Feature Selection:

- Feature selection is crucial for building an effective model. You want to choose the most relevant features to predict house prices. There are various methods for feature selection, such as:
 - Correlation analysis: Identify features that have a strong correlation with the target variable (e.g., using a correlation matrix).
 - Recursive Feature Elimination (RFE): Use techniques like RFE to iteratively remove the least important features.
 - Feature importance from tree-based models: If you plan to use decision tree-based models (e.g., Random Forest), you can use feature importance scores.

3. Data Splitting:

- Split the dataset into a training set and a testing set (e.g., 80% for training and 20% for testing) to evaluate the model's performance.

4. Model Selection:

- Choose a machine learning model suitable for regression tasks. Some common choices include Linear Regression, Decision Trees, Random Forest, Support Vector Machines, and Gradient Boosting models (e.g., XGBoost).

5. Model Training:

- Fit the selected model to the training data using the features you've chosen.
- Tune hyperparameters, if necessary, using techniques like cross-validation or grid search.

6. Model Evaluation:

- Use the testing dataset to evaluate your model. Common evaluation metrics for regression problems include:
 - Mean Absolute Error (MAE): The average absolute difference between predicted and actual prices.

- Mean Squared Error (MSE): The average of the squared differences between predicted and actual prices.
- Root Mean Squared Error (RMSE): The square root of MSE, providing a measure in the original unit of the target variable.
- R-squared (R2) score: A measure of how well the model explains the variance in the target variable.

7. Visualization and Interpretation:

- Visualize the model's predictions against the actual prices to understand how well it performs. You can use scatter plots or residual plots for this purpose.
- Interpret the model's coefficients or feature importances to understand the impact of each feature on house prices.

8. Fine-Tuning and Iteration:

- Based on the evaluation results and interpretation, you may need to make adjustments to your model. This could involve adding or removing features, changing the model, or fine-tuning hyperparameters.

9. Deployment:

- Once you are satisfied with the model's performance, you can deploy it in a real-world application where it can make predictions on new, unseen data.

```
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

# Split the data

X_train, X_test, y_train, y_test = train_test_split(features, target,
                                                    test_size=0.2, random_state=42)

# Create and train the model

model = LinearRegression()

model.fit(X_train, y_train)
```

```
# Make predictions

y_pred = model.predict(X_test)

# Evaluate the model

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)

print("R-squared:", r2)
```

10. Monitoring and Maintenance:

- Continuously monitor the model's performance and update it as necessary to ensure it remains accurate and relevant.

SUBMITTED BY,

STUDENT REG NO: 711221104021

NAAN MUDHALVAN: au711221104021